



PCA for Population Genetics

Nikolay Oskolkov, Group Leader (PI) at NIRI, 28.09.2026



PCA Introduction

Coding in R:

```
data_centered <- scale(data, center = TRUE, scale = FALSE)
```

```
covariance <- t(data_centered) %*% data_centered
```

```
eig <- eigen(covariance)
```

```
plot(eig$vectors[,1:2]);
```



```
barplot(eig$values / sum(eig$values))
```

Mathematically:

$$M_{ij} = X_{ij} - \mu_j$$

$$A = (1/N) * M^T * M$$

$$A * u = \lambda * u$$




It can be analytically derived that the eigen value decomposition in PCA is equivalent to projecting data on axes of maximal variation in the data

Toy example of PCA for population genetics


```
1 gen <- t(matrix(c(1, 0, 2, 0, 2, 0, 2, 1, 1, 1, 0, 1, 0, 2, 1, 2, 1, 1, 1, 1, 1,
2 0, 1, 0, 2, 0, 1, 1, 0, 2, 1, 2, 0, 1, 0), 5, by = TRUE))
3 colnames(gen) <- paste0("Ind", 1:5)
4 rownames(gen) <- paste0("SNP", 1:7)
5 print(gen)
```

	Ind1	Ind2	Ind3	Ind4	Ind5
SNP1	1	1	1	0	0
SNP2	0	1	2	1	2
SNP3	2	1	1	0	1
SNP4	0	0	1	2	2
SNP5	2	1	1	0	0
SNP6	0	0	1	1	1
SNP7	2	2	1	1	0

```
1 gen_centered <- scale(gen, center = TRUE, scale = FALSE)
2 covariance <- t(gen_centered) %*% gen_centered
3 eig <- eigen(covariance)
4 barplot(eig$values/sum(eig$values), names = paste0("PC", 1:5))
```




```
1 plot(eig$vectors[, 1:2], xlab = "PC1", ylab = "PC2", pch = 16, cex = 5, col = 2:6)
2 points(eig$vectors[, 1:2], pch = as.character(1:5))
```



Idea: project the data into a low dimensional space that explains the largest amount of variance

How maximization of covariance is related to eigen value decomposition

- Collapse p features ($p \gg n$) to few latent features and keep variation
- Rotation and shift of coordinate system toward maximal variance
- PCA is an **eigen matrix decomposition** problem



$$PC = u^T X = X^T u$$

X is mean centered $\implies \langle PC \rangle = 0$

$$\langle (PC - \langle PC \rangle)^2 \rangle = \langle PC^2 \rangle = u^T X X^T u$$

$$X X^T = A$$


$$\langle PC^2 \rangle = u^T A u$$

A is **variance-covariance** of X

$$\max(u^T A u + \lambda(1 - u^T u)) = 0$$


$$A u = \lambda u$$

$$X_{ij} \approx U_{ik} V_{kj}$$




Data

$\approx$




Low-dimensional
data representation
(embeddings)



Loadings

$$\text{Loss} = \sum_{i=1}^N \sum_{j=1}^P (X_{ij} - U_{ik} V_{kj})^2$$



Population Structure and Eigenanalysis

Nick Patterson^{1*}, Alkes L. Price^{1,2}, David Reich^{1,2}

¹ Broad Institute of Harvard and MIT, Cambridge, Massachusetts, United States of America, ² Department of Genetics, Harvard Medical School, Boston, Massachusetts, United States of America

Current methods for inferring population structure from genetic data do not provide formal significance tests for population differentiation. We discuss an approach to studying population structure (principal components analysis) that was first applied to genetic data by Cavalli-Sforza and colleagues. We place the method on a solid statistical footing, using results from modern statistics to develop formal significance tests. We also uncover a general “phase change” phenomenon about the ability to detect structure in genetic data, which emerges from the statistical theory we use, and has an important implication for the ability to discover structure in genetic data: for a fixed but large dataset size, divergence between two populations (as measured, for example, by a statistic like F_{ST}) below a threshold is essentially undetectable, but a little above threshold, detection will be easy. This means that we can predict the dataset size needed to detect structure.

Citation: Patterson N, Price AL, Reich D (2006) Population structure and eigenanalysis. *PLoS Genet* 2(12): e190. doi:10.1371/journal.pgen.0020190

Introduction

A central challenge in analyzing any genetic dataset is to explore whether there is any evidence that the samples in the data are from a population that is structured. Are the individuals from a homogeneous population or from a population containing subgroups that are genetically distinct? Can we find evidence for substructure in the data, and quantify it?

This question of detecting and quantifying structure arises in medical genetics, for instance, in case-control studies where uncorrected population structure can induce false positives [1]. It also arises in population genetics, where understanding of the structure may be important to the key scientific issues, especially uncovering the demographic history of the population under study.

We focus on principal components analysis (PCA), which was first introduced to the study of genetic data almost thirty

practical, even on the largest datasets. This is our main aim in this paper.

Using some recent results in theoretical statistics, we introduce a formal test statistic for population structure. We also discuss testing for *additional* structure after some structure has been found. Finally, we are able to estimate the degree of population differentiation that will be detectable for a given data size.

Our methods work in a broad range of contexts, and can be modified to work with markers in linkage disequilibrium (LD). The methods are also able to find structure in admixed populations such as African Americans—that is, in which individuals inherit ancestry from multiple ancestral populations—as long as the individuals being studied have different proportional contributions from the ancestral populations.

We believe that principal components methods largely fell out of favor with the introduction of the sophisticated cluster-based program STRUCTURE [9,10]. STRUCTURE

One concern with applying this approach to genetic data is that the entries in the matrix M do not have the Gaussian distributions expected for a Wishart matrix; instead, they correspond to the three possible genotypes at each SNP. However it is not critical that the entries in the $m \times n$ matrix M be Gaussian. Soshnikov [25] showed that the same TW limit arose if the cell entries were any distribution with high-order moments no greater than the Gaussian. The matrix X is a sum of n rank 1 matrices, and Soshnikov’s result suggests that the same limit would be obtained from any probability distribution in which the columns of M are independent, isotropic (all directions are equiprobable), and such that the column norms have moments no larger than those for a column of independent Gaussian entries. In all our genetic applications, the column norms are in fact bounded, so we can expect the sample covariance matrices to behave well. This theory, originally developed for the case of Gaussian matrix entries, thus seemed likely to work well with large genetic biallelic data arrays. The remainder of this paper verifies that this is the case.

Both PCA and FA optimize a Euclidean distance (L2 norm) between the original data and its low-rank approximation. L2 norm is a consequence of data normality assumption. Why do people still use PCA / FA for non-normally distributed categorical data such as genetic variation?

This norm is often used for convenience because it allows easy and quick/immediate optimization of the problem. So it is true that the loss/ or the model could be changed. There is a trade off between the "quality/accuracy/reality" of the model and the optimization procedure. What we personally experimented with, rather than changing the loss was to add information in the form of constraints on the U or V matrices (e.g. Caye et al <https://hal.archives-ouvertes.fr/hal-01676712>). It works but leads to more complex non deterministic optimization procedure. For other models you can think about Dystruct, Admixture, treemix etc.

Toy example of MDS for population genetics

```
1 gen <- t(matrix(c(1, 0, 2, 0, 2, 0, 2, 1, 1, 1, 0, 1, 0, 2, 1, 2, 1, 1, 1, 1, 1,
2 0, 1, 0, 2, 0, 1, 1, 0, 2, 1, 2, 0, 1, 0), 5, by = TRUE))
3 colnames(gen) <- paste0("Ind", 1:5)
4 rownames(gen) <- paste0("SNP", 1:7)
5 print(gen)
```

	Ind1	Ind2	Ind3	Ind4	Ind5
SNP1	1	1	1	0	0
SNP2	0	1	2	1	2
SNP3	2	1	1	0	1
SNP4	0	0	1	2	2
SNP5	2	1	1	0	0
SNP6	0	0	1	1	1
SNP7	2	2	1	1	0


```
1 my_dist <- dist(t(gen), method = "manhattan", upper = TRUE, diag = TRUE)
2 my_dist
```

	Ind1	Ind2	Ind3	Ind4	Ind5
Ind1	0	3	7	10	11
Ind2	3	0	4	7	8
Ind3	7	4	0	5	4
Ind4	10	7	5	0	3
Ind5	11	8	4	3	0

```
1 my_mds <- cmdscale(my_dist)
2 my_mds
```

	[,1]	[,2]
Ind1	-6.114866	0.2977272
Ind2	-3.077094	0.2402932
Ind3	0.621230	-1.7944048
Ind4	3.723118	2.1483083
Ind5	4.847612	-0.8919238





```
1 plot(my_mds, xlab = "MDS1", ylab = "MDS2", pch = 16, cex = 5, col = 6:2)
2 points(my_mds, pch = as.character(5:1))
```



Idea: project the data into a low dimensional space that preserves distances

Significant Principal Components

- Largest eigen values follow Tracy-Widom distribution
- Therefore we can implement a formal statistical test for detecting population structure in genetic variation data
- PCA dimensionality reduction is directly related to cluster analysis
- Number of largest PCA eigenvalues is approximately the number of clusters
- There is some sort of “phase transition” when population structure becomes visible



COMMENTARY



doi:10.1111/evo.13835



Be careful with your principal components

Mats Björklund^{1,2}

¹Department of Animal Ecology, Evolutionary Biology Centre, Uppsala University, Uppsala, Sweden

²E-mail: mats.bjorklund@ebc.uu.se

Received February 27, 2019

Accepted July 24, 2019

Principal components analysis (PCA) is a common method to summarize a larger set of correlated variables into a smaller and more easily interpretable axes of variation. However, the different components need to be distinct from each other to be interpretable otherwise they only represent random directions. This is a fundamental assumption of PCA and, thus, needs to be tested every time. Sample correlation matrices will always result in a pattern of decreasing eigenvalues even if there is no structure. Tests are, therefore, needed to discern real patterns from illusionary ones. Furthermore, the loadings of the vectors need to be larger than expected by random data to be useful in the calculation of PC-scores. PC-scores calculated from nondistinct PC's have very large standard errors and cannot be used for biological interpretations. I give a number of examples to illustrate the potential problems with PCA. Robustness of the PC's increases with increasing sample size but not with the number of traits. I review a few simple test statistics appropriate for testing PC's and use a real-world example to illustrate how this can be done using randomization tests. PCA can be very useful but great care is needed to avoid spurious results.

KEY WORDS: Correlations, principal components analysis, randomization, standard error.

A common problem in evolutionary biology is that a researcher has a large number of measurements of various variables that might be interpreted for the understanding of, for example, ecological patterns or multivariate selection. One approach to handle this is to simply test variables one by one. However, these variables can be correlated to each other to various degrees, which means that the tests are not independent and the researcher can drown in a multitude of tests that are not easy to interpret.


One very common solution is to try to reduce the dimensionality of the data into fewer independent variables that take the interdependence of variables into account, for example, by a principal component analysis (PCA). One example is given in Figure 1 where four variables, let us say morphological measurements, are correlated to various degrees (the number of traits is arbitrarily chosen to for illustration only). The PCA shows that the first prin-

of the products of the loadings and the individual values creating individual PC-scores): In this example this is a size vector and consequently only one correlation is made using the PC-scores. In this case, it can be inferred that size is negatively correlated with the ecological variable. This is a standard procedure to handle complex data in a more easily understood way. This makes a lot of sense and is not a problem per se.

THE PROBLEM

The problem with the example in Figure 1 is that the correlation matrix is a sample correlation matrix based on a sample size of $N = 20$, while the true correlations are all set to 0.0 (for details see Appendix 1). Thus, there is no pattern apart from that created from sampling error alone. Indeed, in all cases where something is estimated, the estimated value will deviate from the true value

A







Table 1. Statistics from HGDP African Data

Number	Eigenvalue	TW Statistic	TW p -Value	ANOVA p -Value
1	2.07	46.2	$<10^{-12}$	$<10^{-12}$
2	1.40	6.717	3.08×10^{-7}	$<10^{-12}$
3	1.31	0.380	.108	.74



of population structure, and the missing data concentrated in particular samples, populations, and SNPs. For this reason, we only used markers for analysis for which there was no missing data, and we corrected for LD using our regression technique (see below). The details of the data preprocessing

Table 2. Statistics from Thai/Chinese/Japanese Data

Number	Eigenvalue	TW Statistic	TW p -Value	ANOVA p -Value
1	2.21	92.34	$<10^{-12}$	$<10^{-12}$
2	1.47	31.15	$<10^{-12}$	$<10^{-12}$
3	1.23	-1.61	.61	.97

Table 3. Statistics from Shriver Dataset

Number	Eigenvalue	TW Statistic	TW p -Value	ANOVA p -Value
1	22.36	76.091	$<10^{-12}$	$<10^{-12}$
2	8.20	106.870	$<10^{-12}$	$<10^{-12}$
3	5.09	106.071	$<10^{-12}$	$<10^{-12}$
4	3.81	103.146	$<10^{-12}$	$<10^{-12}$
5	3.33	115.239	$<10^{-12}$	$<10^{-12}$
6	2.09	60.090	$<10^{-12}$	$<10^{-12}$
7	1.89	51.768	$<10^{-12}$	$<10^{-12}$
8	1.44	14.658	$<10^{-12}$	$<10^{-12}$
9	1.30	2.038	.010	1.09×10^{-7}
10	1.27	0.084	.084	0.78

A third issue is that our methods require that divergence is small, and that allele frequencies are divergent primarily because of drift. We attempted to apply our methods to an African-American dataset genotyped on a panel of ancestry-informative markers [34]. The Tracy–Widom theory breaks down here with dozens of “significant” axes.





Figure 6. The BBP Phase Change

We ran a series of simulations, varying the sample size m and number of markers n but keeping the product at $mn = 2^{20}$. Thus the predicted phase change threshold is $F_{ST} = 2^{-10}$. We vary F_S and plot the log p -value of the Tracy–Widom statistic. (We clipped $-\log_{10} p$ at 20.) Note that below the threshold there is no statistical significance, while above threshold, we tend to get enormous significance.

doi:10.1371/journal.pgen.0020190.g006

PCA in Population Genetics: Literature Review



SHARE **ARTICLES**

Synthetic maps of human gene frequencies in Europeans



P Menozzi, A Piazza, L Cavalli-Sforza
 See all authors and affiliations

Science 01 Sep 1978;
 Vol. 201, Issue 4358, pp. 786-792
 DOI: 10.1126/science.356262

Article Info & Metrics eLetters PDF

Abstract

Multivariate techniques can be used to condense the information for a large number of loci and alleles into one or a few synthetic variables. The geographic distribution of synthetic variables can be plotted by the same technique used in mapping the gene frequency of a single allele. Synthetic maps were constructed for Europe and the Near East, with the use of principal components to condense the information of 38 independent alleles from ten loci. The first principal component summarizes close to 30% of the total information and shows gradients. Maps thus constructed show clines in remarkable agreement with those expected on the basis of the spread of early farming in Europe, thus supporting the hypothesis that this spread was a demic spread rather than a cultural diffusion of farming technology.



Genetic distances correlate with geographical distances possibly due to people's migration

Neolithic farming migration rather than a cultural diffusion of farming technology

Principal component analysis of genetic data

David Reich, Alkes L Price & Nick Patterson

Principal component analysis (PCA) has been a useful tool for analysis of genetic data, particularly in studies of human migration. A new study finds evidence that the observed geographic gradients, traditionally thought to represent major historical migrations, may in fact have other interpretations.

Principal component analysis (PCA) has been used for several decades to study human population migrations, resulting in remarkable inferences about history. On page 646 of this issue, John Novembre and Matthew Stephens¹ show that the geographic gradients that emerge when PCA is applied to genetic data—and that are sometimes interpreted as highly suggestive of major historical migrations—can also have other explanations. We suggest guidelines for scientists interested in using PCA in genetic analysis in light of this potential concern and highlight three applications in which PCA has continued value: detecting population substructure, correcting for stratification in disease studies and making qualified inferences about human history.

Synthetic maps in question

PCA is a statistical method for exploring and making sense of datasets with a large number of measurements (which can be thought of as dimensions) by reducing the dimensions to the few principal components (PCs) that explain the main patterns. Thus, the first PC is the mathematical combination of measurements that accounts for the largest amount of variability in the data. Luca Cavalli-Sforza and colleagues had the original insight that PCA could be applied to human genetic variation², and they eventually analyzed




Figure 1 PCA continues to provide evidence of important migration events. (a) We carried out PCA on 940 individuals from the Human Genome Diversity Project that were scanned at approximately 650,000 SNPs^{1,1} using data from 101 sub-Saharan African samples to define the PCs (Mandenka, Bantu from Kenya and South Africa, Yoruba, San, Mbui Pygmy and Biaka Pygmy). We carried out the analysis on samples blinded to population labels (the coloring of samples was only carried out after the analysis). We plotted principal component 1 (negative values are more Bantu-related) and principal component 4 (positive values are more closely related to the Senegalese Mandenka). (b) Outlying populations are the Mozabite, who are more Mandenka-related, reflecting recent gene flow across the Sahara, and Papuans and Melanesians, who have inherited less Bantu-related gene flow. (c) To reveal the west-to-east gradient of Bantu-related ancestry across Eurasia, we averaged the first PC for each of the non-African populations and plotted the populations in rank order.

Interpreting principal component analyses of spatial population genetic variation

John Novembre^{1,3} & Matthew Stephens^{1,2}



Nearly 30 years ago, Cavalli-Sforza *et al.* pioneered the use of principal component analysis (PCA) in population genetics and used PCA to produce maps summarizing human genetic variation across continental regions¹. They interpreted gradient and wave patterns in these maps as signatures of specific migration events^{1–3}. These interpretations have been controversial^{4–7}, but influential⁸, and the use of PCA has become widespread in analysis of population genetics data^{9–13}. However, the behavior of PCA for genetic data showing continuous spatial variation, such as might exist within human continental groups, has been less well characterized. Here, we find that gradients and waves observed in Cavalli-Sforza *et al.*'s maps resemble sinusoidal mathematical artifacts that arise generally when PCA is applied to spatial data, implying that the patterns do not necessarily reflect specific migration events. Our findings aid interpretation of PCA results and suggest how PCA can help correct for continuous population structure in association studies.

Cavalli-Sforza *et al.*'s classic text “The History and Geography of Human Genes”¹⁴ synthesizes a decades-long survey of human genetic variation. These ground-breaking datasets stimulated development of methods that are now widely used, including application of principal component analysis (PCA) to population genetic variation. In essence, Cavalli-Sforza *et al.* collected count data for many genetic variants




origin must be the center of the radiation” (p. 295 of ref. 3). They also suggest centripetal population movements as an alternative explanation. Examples of their explanations for the European PC maps in Figure 1 include expansion of agriculturalists out of the Near East (Europe PC1); migrations of Mongoloid Uralic speakers from north-western Asia (Europe PC2); migration of the carriers of the proto-Indo-European Kurgan culture in Europe (Europe PC3); and an expansion from Greece (Europe PC4).


Because the basis for these interpretive guidelines is unclear, we performed simulations to investigate whether such specific migration events are necessary to explain the observed patterns. Specifically, we performed PCA on data simulated under equilibrium population genetic models without range expansions, assuming a constant homogeneous short-range migration process across both time and (two-dimensional) space. The results showed highly distinctive structure. For example, the first two PC maps show large-scale orthogonal gradients, and the next two show ‘saddle’ and ‘mound’ patterns (Fig. 1). The same four basic patterns occurred consistently in the first few PC maps across multiple simulations, although not always in the same order (Supplementary Fig. 1 online). Results for the analogous one-dimensional habitat setting are even more structured, resembling sinusoidal functions of increasing frequency (Fig. 2, Supplementary Fig. 2 online). Thus PC maps show local peaks and troughs even when underlying migration patterns are homogeneous across time and space. This suggests that the local features of the PC

Genetic distances correlate with geographical distances possibly due to isolation by distance and not necessarily migration




■ French-speaking Swiss ■ French
■ German-speaking Swiss ■ German
■ Italian-speaking Swiss ■ Italian








Novembre & Stephens, Nat. Genetics 2008



Francois & Jay, Nat. Communications 2020






“The results described here provide an explanation. First, from Equation 10 it can be seen that the the **matrix M is influenced by the relative sample size from each population** through the components t_i . For instance, even if all populations are equally divergent from each other, those for which there are fewer samples will have larger values of t_i because relatively more pairwise comparisons are between populations”

$$A = XX^T = \frac{1}{N} \sum_{ij} x_i x_j$$


$$N = N_{pop1} + N_{pop2} + N_{pop3} + \dots = \sum_k N_k$$

$$A_{uneven} = \sum_{ijk} \frac{1}{N_k} x_{ik} x_{jk}$$


Figure 3. The effect of uneven sampling on PCA projection. PCA projection of samples taken from a set of nine populations arranged in a lattice, each of which exchanges migrants at rate M per N_e generations with each adjoining neighbour, leads to a recovery of the migration-space if samples are of equal size (A), or a distortion of migration-space if populations are not equally represented (B,C). In each part the left-hand panel shows the analytical solution (the area of each point represents the relative sample size) with migration routes illustrated while the right-hand panel shows the result of a simulation with a total sample size of 180 and 10,000 independent SNP loci. All examples are for $M = 2$.
doi:10.1371/journal.pgen.1000686.g003



Downsampled Europeans



Downsampled Asians



Correcting For Unequal Sampling in Principal Component Analysis of Genetic Data

W. Burgos-Paz^{1,2}, S. E. Ramos-Onsins¹, M. Pérez-Enciso^{1,2,3}, L. Ferretti⁴

¹ Centre for Research in Agricultural Genomics (CRAG), 08193 Bellaterra, Spain, ² Departament de Ciència Animal i dels Aliments, Universitat Autònoma de Barcelona, 08193, Bellaterra, Spain, ³ Institut Català de Recerca i Estudis Avançats (ICREA), Carrer de Lluís Companys 23, Barcelona, 08010, Spain, ⁴ UMR 7138, UPMC and CIRB, College de France, Paris, France

ABSTRACT: Principal component analysis (PCA) is one of the most widely used tools to explore variability of high dimensional data. PCA is used for population and quantitative genetics. Its popularity has recently increased due to the huge amount of molecular markers available in datasets worldwide. In genetics, a common issue due to external constraints is uneven sampling of populations, limiting the usefulness of PCA because of well-known sample size sensitivity and two-dimensional projection bias. Here we evaluated the use of weighted PCA (wPCA) in genetic data in order to correct uneven sampling bias. Simulations suggest that wPCA improves the two-dimensional projections of PCA data and, in some cases, recovers population relationships patterns, even when sample size is as low as $n=1$. We used this correction in pig data from populations with uneven sampling, recovering a more realistic structure than inferred with only PCA.

Keywords: SNP; population structure; phylogeography

Introduction

The Principal Component Analysis (PCA) is a widely used tool for visualization of the structure of a highly multidimensional data set. It has applications in many fields, including genetics (Novembre and Stephens (2008)). The idea behind this technique is to reduce the complexity of the data by retaining only the linear combinations of variables explaining the maximum variance in the data set. Individual data is then projected onto the space of these combinations and visualized as a low-dimensional plot. In population genetics, PCA is frequently applied to visualize the genetic structure of populations, and its popularity has increased in the last years because of the high throughput genotyping technologies (e.g., SNP arrays), where it is necessary to extract the underlying structure of the dataset

McVean (2009) suggested to correct for unequal sampling by downsampling the different populations, but this may lead to a loss in power or in a less accurate picture of genetic structure. A natural framework for this is to correct for this issue by means of a weighted PCA (wPCA), where the variables measured for each sample can be assigned a different weight w_i (Kriegel et al., (2008)). Kriegel et al. proposed this correction for synthetic data to evaluate the influence of outliers on PCA but, to our knowledge, its properties in genetic data are unknown. Here, we evaluated the usefulness of this correction with data from SNPs polymorphism under different population relationships, and illustrate the method in pig populations where differences in sample size could under-over estimate population relationships.

Materials and Methods

Weighted PCA. Given a set of genetic data of n sequences/genotypes and s biallelic SNPs for each sequence, the first step of the PCA is to compute the covariance:

$$\begin{aligned} Cov(x_s, x_{s'}) &= E(x_s, x_{s'}) - E(x_s)E(x_{s'}) \\ &= \frac{1}{n} \sum_{ij} x_{ij}^s x_{ij}^{s'} - \left(\frac{1}{n} \sum_i x_i^s \right) \left(\frac{1}{n} \sum_i x_i^{s'} \right) \end{aligned}$$

A common assumption is that all individuals should be equally important in the PCA (that is, the implicit weight of each individual is $1/n$). This is usually correct, since there is often no prior information about the origin and grouping of individuals. However, this is not true for samples taken from different breeds or structured natural populations living in different habitats or geographical locations, since the degree of relationship will differ

manner (Paschou et al., (2007)).

Despite several advantages of PCA in the population genetics field, this technique is very sensitive to the choice of the dataset and the distortion of the PCA plots due to a biased or unequal sampling is a known problem (McVean (2009)). In genetic studies, it is often not possible to sample individuals according to some criteria chosen a priori. For instance, the choice of sampling could be biased by incomplete knowledge of the population structure, conservation or there could also be ethical issues. Further, many factors, like budget constraints or accessibility of geographical regions, availability of samples and technical problems in their conservation and sequencing limit sampling in the field.

$$w_i = \frac{1}{n_{pop}} \cdot \frac{1}{n_A}, \quad i \in A$$

Simulations. To evaluate the usefulness of wPCA in genetic data, we simulated six populations considering two scenarios where population relationships are mainly shaped by the individuals exchange (e.g., migrations or introgression). In the first scenario (M1), we considered that all populations are equally related in terms of exchange of individuals. This scenario may occur in cooperative breeding schemes where sires or dams are regularly interchanged between herds. In the second scenario (M2), a subpopulation structure was considered where some populations have high individual exchange rates among them but low with the rest. This will typically occur, e.g., when one nucleus provides genetic material to other herds that are otherwise not connected.

For each population, five hundred individuals with thousand independent SNPs each were simulated by coalescence using MLCOASIM.v2 (available at <http://bioinformatics.cragenomica.es/numgenomics/people/sbas/software/software.html>). Effective population size ($N_e=1000$) was similar for all populations. Further, each population was sampled retaining one hundred individuals each for downstream analyses. To simulate unbalanced

In the wPCA paradigm, each sample's genotype can be assigned a different weight w_i ; then, the covariance can be rewritten naturally as:

$$Cov(x_s, x_{s'}) = \sum_i w_i x_{ij}^s x_{ij}^{s'} - \left(\sum_i w_i x_i^s \right) \left(\sum_i w_i x_i^{s'} \right),$$

where the weights are numbers between $0 < w_i < 1$ with $\sum_i w_i = 1$. This reduces to the usual PCA when $w_i = 1/n$. We denote the number of populations by n_{pop} and the number of sequences in population A by n_A . Since the weight of population A is the sum of the weights of all the sequences belonging to A , a natural weight for each sequence of A is:

the PCA projection with an unbalanced finite dataset using the usual (middle) and weighted (right) approach. An unweighted PCA result in highly distorted PC projections, in agreement with McVean, (2009). The PCA distortion was especially noticeable in M1, where individuals from reduced populations were located in the middle of the populations with larger sample size. In M2, distortion may cause misinterpretation of subpopulation structure (i.e blue dot individuals).




Figure 1. PCA and wPCA projections of simulated

1000G: PCA covariance matrix corrected for uneven population size


$$A = XX^T = \frac{1}{N} \sum_{ij} x_i x_j$$

$$N = N_{pop1} + N_{pop2} + N_{pop3} + \dots = \sum_k N_k$$

$$A_{uneven} = \sum_{ijk} \frac{1}{N_k} x_{ik} x_{jk}$$

Normalize each sample by its population size before computing the covariance matrix

Is it still an unsupervised technique?














*Knut och Alice
Wallenbergs
Stiftelse*





LUNDS
UNIVERSITET







← scRNAseq






← Genetic Variation





1000G_EUR: EFFECT OF NUMBER OF SAMPLES ON VARIANCE EXPLAINED













In other OMICs the observed variance explained is typically well separated from the variance explained of the shuffled data matrix




Pairwise distances grow as:

$$\text{dist} \sim \sqrt{p}$$

data points become far apart

Dep. Variable:	y	R-squared (uncentered):	1.000			
Model:	OLS	Adj. R-squared (uncentered):	1.000			
Method:	Least Squares	F-statistic:	1.454e+06			
Date:	Mon, 01 Feb 2021	Prob (F-statistic):	0.00			
Time:	17:48:02	Log-Likelihood:	734.17			
No. Observations:	498	AIC:	-1466			
Df Residuals:	497	BIC:	-1462			
Df Model:	1					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
x1	0.5636	0.000	1205.756	0.000	0.563	0.565
Omnibus:	100.625	Durbin-Watson:	0.007			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	161.978			
Skew:	1.274	Prob(JB):	6.71e-36			
Kurtosis:	4.145	Cond. No.	1.00			



MNIST
benchmark dataset

0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
 3 3 3 3 3 3 3 3 3 3 3 3 3 3 3
 4 4 4 4 4 4 4 4 4 4 4 4 4 4 4
 5 5 5 5 5 5 5 5 5 5 5 5 5 5 5
 6 6 6 6 6 6 6 6 6 6 6 6 6 6 6
 7 7 7 7 7 7 7 7 7 7 7 7 7 7 7
 8 8 8 8 8 8 8 8 8 8 8 8 8 8 8
 9 9 9 9 9 9 9 9 9 9 9 9 9 9 9

It looks like replacing original data by UMAP components can mitigate the harm from the Curse of Dimensionality and allow better clustering